

Project Report: Transformer-Based Trading Model

Model Selection:

Pretrained Transformer (**BERT**): I started with a pretrained BERT model for sequence classification. The model is fine-tuned in provided dataset to predict the trading signals. BERT was chosen for its ability to handle classification tasks with strong performance.

Feature Engineering(Used Existing momentum provided in a PPO Code)

1. Selection of Technical Indicators:

- Close Price (**Close**): The closing price of the asset.
- Relative Strength Index (**RSI**): A momentum indicator that measures the magnitude of recent price changes.
- Moving Average Convergence Divergence (**MACD**): A trend-following momentum indicator.
- Stochastic Oscillator K (**Stoch_k**): A momentum indicator comparing a particular closing price to a range of its prices over time.
- Average True Range (**ATR**): Used to measure volatility.
- Other Indicators: Indicators such as **OBV**, Bollinger Bands, ADX, and others were also considered to capture different market conditions.

2. Textual Representation of Features:

The indicators were combined into a textual format that could be tokenized and fed into the transformer model. Example:

```
text = (  
    f"Price is {row['Close']}, RSI is {row['RSI']}, MACD is {row['MACD']}, "  
    f"Stochastic K is {row['Stoch_k']}, ATR(14) is {row['ATR_14']}, "  
    f"ADX is {row['ADX']}"  
)
```

Hyperparameter Tuning:

1. Initial Setup:

- **Learning Rate:** Started with a standard learning rate of `2e-5`, typical for fine-tuning transformer models.
- **Batch Size:** Set to `16` to balance between computational efficiency and model convergence.
- **Epochs:** Initially set to `5` epochs to observe the model's learning curve, with early stopping configured to prevent overfitting.

2. Learning Rate Scheduling(cosine and linear):

Cosine

```
training_args = TrainingArguments(  
    output_dir='./results',  
    num_train_epochs=5,  
    per_device_train_batch_size=16,  
    per_device_eval_batch_size=16,  
    warmup_steps=500,  
    weight_decay=0.01,  
    logging_dir='./logs',  
    logging_steps=10,  
    eval_strategy="epoch",  
    save_strategy="epoch",  
    load_best_model_at_end=True,  
    learning_rate=2e-5,  
    lr_scheduler_type="cosine",  
)  
  
trainer = Trainer(  
    model=model,  
    args=training_args,  
    train_dataset=dataset,  
    eval_dataset=dataset,  
    callbacks=[EarlyStoppingCallback(early_stopping_patience=3)],  
)  
  
trainer.train()
```

```
Step: 1  
Balance: 38401.62999998344  
Shares held: 51453.0  
Total shares traded: 51453.0  
Total portfolio value: 10022856.279999984  
Cumulative reward: -10857.513904298383
```

Linear

```
training_args = TrainingArguments(  
    output_dir='./results',  
    num_train_epochs=5,  
    per_device_train_batch_size=16,  
    per_device_eval_batch_size=16,  
    warmup_steps=500,  
    weight_decay=0.01,  
    logging_dir='./logs',  
    logging_steps=10,  
    eval_strategy="epoch",  
    save_strategy="epoch",  
    load_best_model_at_end=True,  
    learning_rate=2e-5,  
    lr_scheduler_type="linear",  
)  
  
trainer = Trainer(  
    model=model,  
    args=training_args,  
    train_dataset=dataset,  
    eval_dataset=dataset,  
    callbacks=[EarlyStoppingCallback(early_stopping_patience=3)],  
)  
  
trainer.train()
```

```
Step: 1  
Balance: 3610311.9199999445  
Shares held: 33164.0  
Total shares traded: 33164.0  
Total portfolio value: 10045786.119999945  
Cumulative reward: -11746.744797677315
```

- Implemented a linear learning rate scheduler with warmup to gradually increase the learning rate at the beginning of training and then decrease it, helping the model converge more smoothly.
- Warmup Steps: Initially set to `500` to allow the model to start with smaller updates.

4. Early Stopping:

- Early Stopping: Configured with a patience of `3`, monitoring the validation loss to halt training if no improvement was observed, thus saving time and resources.

Integrate with PPO Model:

- **Environment Setup:** The trading environment (**TradingEnvironmentWithTransformer**) was customized to incorporate both the Transformer model's predictions and the PPO model's actions.
- **_take_action Method:**
 - Transformer model was used to predict the trading signal (Buy, Sell, Hold).
 - The PPO model received these predictions and other state variables as inputs to make a final decision.

Model Performance:

- Training Loss: Final training loss settled at 0.240900.
- Validation Loss: Final validation loss improved to 0.225595, indicating good generalization performance.